



SDU Summer School

# Deep Learning

Fall 2022

## Introduction to KERAS

# What is Keras?



- Keras is a high-level API providing easy to use elements for deep learning
- Can work with several backends
- Programs can easily be deployed on CPUs, GPUs without changing the code

# Who makes Keras? Contributors and backers

 633 contributors

Google



# The Keras user experience

- Keras API is easy to understand.
  - Keras follows best practices for reducing cognitive load: it offers consistent & simple APIs, it minimizes the number of user actions required for common use cases, and it provides clear and actionable feedback upon user error.
- Easy to learn.
  - As a Keras user, you are more productive, allowing you to try more ideas than your competition, faster -- which in turn helps you win machine learning competitions.
- This ease of use does not come at the cost of reduced flexibility:
  - Keras integrates with lower-level deep learning languages (in particular TensorFlow), it enables you to implement anything you could have built in the base language. In particular, as `tf.keras`, the Keras API integrates seamlessly with your TensorFlow workflows.

# Multi-Backend, Multi-Platform

- Develop in Python, R
  - On Unix, Windows, OSX
- Run the same code with...
  - TensorFlow
  - Cognitive Toolkit (CNTK) - Microsoft
  - Theano
  - MXNet
  - PlaidML
  - ??
- Run on CPU, NVIDIA GPU, AMD GPU, TPU...



# How to use Keras: An introduction

# Three API Styles

- **The Sequential Model**
  - Dead simple
  - Only for single-input, single-output, sequential layer stacks
  - Good for 70+% of use cases
- **The functional API**
  - Like playing with Lego bricks
  - Multi-input, multi-output, arbitrary static graph topologies
  - Good for 95% of use cases
- **Model subclassing**
  - Maximum flexibility
  - Larger potential error surface

# The Sequential API

```
import tensorflow as tf
from tensorflow.keras import layers

model = tf.keras.Sequential()
model.add(layers.Dense(20, activation='relu', input_shape=(10,)))
model.add(layers.Dense(20, activation='relu'))
model.add(layers.Dense(20, activation='softmax'))

model.fit(x, y, epochs=10, batch_size=32)
```



# Options for Layers

- Core Layers
- Convolutional Layers
- Pooling Layers
- Locally-connected Layers
- Recurrent Layers
- Reshape Layers
- Dropout Layers
- Merge Layers
- Normalization Layers
- Noise layers

# Options for Layers

- The core layers perform the most basic operations
- They are enough to build FFN networks
- Core Layers
  - Input Layers
  - Dense Layers
  - Activation Layer
  - Embedding Layers
  - Masking layers
  - Lambda Layers

# Dense Layer

```
tensorflow.keras.layers.Dense(units, #Number of units in the layer
                                #Standard: use linear output
                                use_bias=True, #Add a bias vector
                                kernel_initializer='glorot_uniform', #How to initialize the weights
                                bias_initializer='zeros', #How the biases
                                kernel_regularizer=None, #For example, apply L2 regularization
                                bias_regularizer=None, #For example, apply L2 regularization
                                activity_regularizer=None, #For example, apply L2 regularization
                                kernel_constraint=None, #For example, non-negative constraint
                                bias_constraint=None #For example, non-negative constraint
                                )
```

# Activation Function

```
model.add(Dense(64))
model.add(Activation('tanh'))
#This is equivalent to:
model.add(Dense(64, activation='tanh'))
```

- Available Activations:
  - softmax
  - elu: (Exponential linear unit.)
    - $x$  if  $x > 0$  and  $\alpha * (\exp(x) - 1)$  if  $x < 0$ .
  - selu: Scaled Exponential Linear Unit
  - softplus
    - $\log(\exp(x) + 1)$
  - relu
    - $\text{relu}(x, \alpha=0.0, \text{max\_value}=\text{None})$
  - sigmoid
  - tanh

# Compiling the Model

```
from tensorflow.keras import optimizers

model.compile(
    optimizer=optimizers.RMSprop(lr=0.001),
    loss='binary_crossentropy',
    metrics=['accuracy']
)
```

- Before training a model, you need to configure the learning process, which is done via the compile method, defining
  - **An optimizer.** This could be the string identifier of an existing optimizer (such as rmsprop or adagrad), or an instance of the Optimizer class
  - **A loss function.** This is the objective that the model will try to minimize. It can be the string identifier of an existing loss function, or it can be an objective function.
  - **A list of metrics.** A metric could be the string identifier of an existing metric or a custom metric function.

# Examples: Compiling Models

```
# For a multi-class classification problem
model.compile(optimizer='adam',
              loss='categorical_crossentropy',
              metrics=['accuracy'])

# For a binary classification problem
model.compile(optimizer='rmsprop',
              loss='binary_crossentropy',
              metrics=['accuracy'])

# For a mean squared error regression problem
from tensorflow.keras import optimizers
model.compile(optimizer=optimizers.RMSprop(lr=0.0025),
              loss='mse')
```

# What does compile do?

- Compile defines the loss function, the optimizer and the metrics. That's all.
  - It has nothing to do with the weights and you can compile a model as many times as you want without causing any problem to pretrained weights.
  - You need a compiled model to train (because training uses the loss function and the optimizer). But it's not necessary to compile a model for predicting.
  - Do you need to use compile more than once? Only if:
    - You want to change one of these:
      - Loss function
      - Optimizer / Learning rate
      - Metrics
    - You loaded (or created) a model that is not compiled yet. Or your load/save method didn't consider the previous compilation.
- **Consequences of compiling again:**
  - If you compile a model again, you will lose the optimizer states.

# Loss Functions

- mean\_squared\_error
- mean\_absolute\_error
- mean\_absolute\_percentage\_error
- mean\_squared\_logarithmic\_error
- **binary\_crossentropy**
- **categorical\_crossentropy**
- sparse\_categorical\_crossentropy
- ...



# Metrics

- Can be any of the loss functions
- Some standard metrics like
  - F1
  - Precision
  - Recall
  - accuracy

# Train the Model

```
model.fit(x=None, y=None,
          batch_size=None,
          epochs=1,
          verbose=1,
          callbacks=None,
          validation_split=0.0,
          validation_data=None,
          shuffle=True,
          class_weight=None,
          sample_weight=None,
          ...
)
```

# Input and desired outcome  
# Number of samples per gradient update. If none, it defaults to 32  
# Number of runs over the complete x and y  
# Verbosity mode. 0 = silent, 1 = progress bar, 2 = one line per epoch.  
# List of functions to call during training  
# Part of dataset set aside for validating  
# Validation dataset, tuple (x\_val, y\_val)  
# shuffle the training data before each epoch  
# Give some classes more/less weight  
# Give some samples more/less weight

# In Context

```
model.compile(optimizer='rmsprop',  
              loss='binary_crossentropy',  
              metrics=['acc'])  
  
history = model.fit(x=partial_x_train, y=partial_y_train,  
                    epochs=20,  
                    batch_size=512,  
                    validation_data=(x_val, y_val))
```

# The History Object

- Note that the call to `model.fit()` returns a History object. This object has a member `history`, which is a dictionary containing data about everything that happened during training.

```
>>> history_dict = history.history
>>> history_dict.keys()
[u'accuracy', u'loss', u'val_accuracy', u'val_loss']
```

- The dictionary contains four entries: one per metric that was being monitored during training and during validation.
- You can now plot these to get Information about your performance

# Plotting the training and validation loss

```
import matplotlib.pyplot as plt


history_dict = history.history
loss_values = history_dict['loss']
val_loss_values = history_dict['val_loss']

epochs = range(1, len(loss_values) + 1)

# 'bo' is for blue dot, 'b' is for solid blue line
plt.plot(epochs, loss_values, 'bo', label='Training loss')
plt.plot(epochs, val_loss_values, 'b', label='Validation loss')
plt.title('Training and validation loss')
plt.xlabel('Epochs')
plt.ylabel('Loss')
plt.legend()

plt.show()
```

# Plotting the training and validation loss



# Practical Recommendations

- For lower data amounts, you should train smaller and shallower networks in order to prevent overfitting
- Preprocessing
  - Take small values - Typically, most values should be in the 0–1 range.
  - Be homogenous- That is, all features should take values in roughly the same range.

# Load and Save models

- You save a Keras model into a single HDF5 file which will contain:
  - the architecture of the model, allowing to re-create the model
  - the weights of the model
  - the training configuration (loss, optimizer)
  - the state of the optimizer, allowing to resume training exactly where you left off.

```
from tensorflow.keras.models import load_model

model.save('my_model.h5') # creates a HDF5 file 'my_model.h5'
del model                 # deletes the existing model

# returns a compiled model
# identical to the previous one
model = load_model('my_model.h5')
```