


SDU Summer School

# Deep Learning

Summer 2022

**Welcome to the Summer School**



# Optimization for Deep Learning

- **Parameter Initialization**
- Batch Normalization
- Pre-Training

# Types of Initialization

1. Non-iterative optimization requires no initialization
  - Simply solve for solution point
2. Iterative but converge regardless of initialization
  - Acceptable solutions in acceptable time
3. Iterative but affected by choice of Initialization
  - Deep learning training algorithms are iterative
  - Initialization determines whether it converges at all
  - Can hugely determine how quickly learning converges

# Modern Initialization Strategies

- They are simple and heuristic
- Based on achieving nice properties
- But problem is a difficult one
  - Some initial points are beneficial for optimization but detrimental to generalization
- Only property known with certainty: Initial parameters must be chosen to **break symmetry**
  - If two hidden units have the same inputs and same activation function, then they must have different initial parameters
  - Usually best to initialize each unit to compute a different function
  - This motivates use random initialization of parameters

# Choice of biases

- Biases for each unit are heuristically chosen constants
- Only the weights are initialized randomly
- Extra parameters such as conditional variance of a prediction are constants like biases

# Weights drawn from Gaussian

- Weights are almost always drawn from a Gaussian or uniform distribution
  - Choice of Gaussian or uniform does not seem to matter much but not studied exhaustively
- Scale of the initial distribution does have an effect on outcome of optimization and ability to generalize
  - Larger initial weights will yield stronger symmetry-breaking effect, helping avoid redundant units
  - Too large may result in exploding values

# Heuristics for initial scale of weights

- One heuristic is to initialize the weights of a fully connected layer with  $N$  inputs and  $M$  outputs by sampling each weights from  $\text{Uniform}(-r, r)$  with

$$r = \frac{1}{\sqrt{N}}$$

- Another heuristic is normalized initiation with

$$r = \sqrt{\frac{6}{N + M}}$$

- Which is a compromise between the goal of initializing all layers to have the same activation variance and the goal of having all layers having the same gradient variance

# Initialization for the biases

- Bias settings must be coordinated with setting weights
- Setting biases to zero is compatible with most weight initialization schemes
- Situations for nonzero biases:
  - Bias for an output unit: initialize to obtain right marginal statistics for output
    - Set bias to inverse of activation function applied to the marginal statistics of the output in the training set
    - Assuming that the weights in the beginning are so small, that output is driven only by biases
  - Choose bias to causing too much saturation at initialization





# Optimization for Deep Learning

- Parameter Initialization
- **Batch Normalization**
- Pre-Training

# Batch Normalization

- Batch normalization: exciting recent innovation
- Method is to replace activations with zero-mean with unit variance activations
- Motivation is difficulty of choosing learning rate  $\epsilon$  in deep networks
  - Method adds an additional step between layers, in which the output of the earlier layer is normalized
    - By standardizing the mean and standard deviation of each individual unit
  - It is a method of adaptive re-parameterization
    - It is not an optimization algorithm at all
    - A method to reduce internal covariate shift in neural networks

# Motivation: Difficulty of composition

- Very deep models involve compositions of several functions or layers

$$f(\mathbf{x}, \mathbf{w}) = f^{(l)} \left( \dots f^{(3)} \left( f^{(2)} \left( f^{(1)}(\mathbf{x}) \right) \right) \right)$$

- The gradient tells us how to update each parameter


$$\mathbf{w}^{\tau+1} = \mathbf{w}^{\tau} - \epsilon \nabla_{\mathbf{w}} J(f(\mathbf{x}, \mathbf{w}), y)$$

- **Under assumption that other layers do not change**
- **BUT:** We update all  $l$  layers simultaneously
- When we make the change unexpected results can happen
- Because many functions are changed simultaneously, effects can accumulate

# Choosing learning rate $\epsilon$ in multilayer

- Simple example:

- $l$  layers, one multiplication unit per layer, no activation function



- Network simply computes

$$\hat{y} = x \cdot w_1 \cdot w_2 \dots \cdot w_l$$

- Output of Layer  $i$  is  $h_i = h_{i-1}w_i$
- Output is a linear function of input  $x$  but a nonlinear function of the weights  $w_i$

# Gradient in Simple example

- Suppose our cost function has put a gradient of 1 on  $\hat{y}$ , so we wish to decrease  $\hat{y}$  slightly.

- The back-propagation algorithm can then compute a gradient

$$\mathbf{g} = \nabla_{\mathbf{w}} \hat{y}$$

- Which corresponds to the gradient  $\mathbf{g}$  evaluated at  $y = \hat{y}$
- When using the update  $\mathbf{w} \leftarrow \mathbf{w} - \epsilon \mathbf{g}$  predicts that  $\hat{y}$  decreases by  $\epsilon \mathbf{g}^T \mathbf{g}$
- If we want to decrease  $\hat{y}$  by 0.1, we could set the learning rate to

$$\epsilon = \frac{0.1}{\mathbf{g}^T \mathbf{g}}$$

# Difficulty of Multilayer learning Rate

- With the first order information, we would set the learning rate to

$$\epsilon = \frac{0.1}{\mathbf{g}^T \mathbf{g}}$$

- Problem: We have to deal with many 2<sup>nd</sup>, 3<sup>rd</sup> ... effects. The new value in fact is

$$x(w_1 - \epsilon g_1)(w_2 - \epsilon g_2) \dots (w_l - \epsilon g_l)$$

- One example for a second order term would be  $\epsilon^2 g_1 g_2 \prod_{i=3}^l w_i$
  - This term might be negligible if  $\prod_{i=3}^l w_i$  is small or exponentially large if weights are larger than 1.
- This makes it very hard to choose  $\epsilon$  because the effects of an update for one layer depend so strongly on all other layers
  - Higher order methods tackle this problem, but still only for small  $l$

# The Batch Normalization Solution

- Provides an elegant way of re-parameterizing almost any network
- Significantly reduces the problem of coordinating updates across many layers
- Can be applied to any input or hidden layer in a network

# Batch Normalization Equations

- **$H$** : minibatch of activations of layer to normalize

- arranged as a design matrix
- With activations for each example appearing in a row

$$H = \begin{matrix} & \begin{matrix} \text{units} & \longrightarrow \end{matrix} \\ \begin{matrix} a_{1,1} & a_{1,2} & a_{1,3} \\ \vdots & \vdots & \vdots \\ a_{N,1} & a_{N,2} & a_{N,3} \end{matrix} & \begin{matrix} \uparrow \\ \text{samples} \end{matrix} \end{matrix}$$

- To normalize  **$H$**  we replace it with  $H' = \frac{H - \mu}{\sigma}$ 
  - Where  $\mu$  is a vector containing the mean of each unit and  $\sigma$  is the std. deviation of each unit
  - The arithmetic here is based on broadcasting the vector  $\mu$  and the vector  $\sigma$  to be applied to every row of  **$H$**
  - Within each row, the arithmetic is element-wise
  - $H_{i,j}$  is normalized by subtracting  $\mu_j$  & dividing by  $\sigma_j$



# Normalization Details

- Rest of the network operates on  $\mathbf{H}'$  in the same way that the original network operated on  $\mathbf{H}$
- At training time

$$\mu = \frac{1}{m} \sum_i H_i \quad \text{and} \quad \sigma = \sqrt{\delta + \frac{1}{m} \sum_i (H - \mu)_i^2}$$

- where  $\delta$  is a small positive value such as  $10^{-8}$  imposed to avoid encountering the undefined gradient of  $z = 0$ 
  - Crucially we back propagate through these operations for computing the mean and std dev
  - And for applying them to normalize  $\mathbf{H}$ 
    - This means that the gradient will never propose an operation that acts simply to increase std dev or mean of  $h_i$  the normalization operations remove the effect of such an action and zero out the component in the gradient

# Batch Normalization at Test time

- At test time,  $\mu$  and  $\sigma$  may be replaced by running averages that were collected during training time
- This allows the model to be evaluated on a single example without needing to use definitions of  $\mu$  and  $\sigma$  that depend on an entire minibatch

# Revisiting the simple example

- Revisiting the  $\hat{y} = x \cdot w_1 \cdot w_2 \dots \cdot w_l$  example we can mostly resolve the difficulties in learning the model by normalizing  $h_{l-1}$
- Suppose that  $x$  is drawn from a unit Gaussian
- Then  $h_{l-1}$  will also come from a Gaussian, because the transformation from  $x$  to  $h_l$  is linear
- However  $h_{l-1}$  will no longer have zero mean, unit variance

# Restoring zero-mean unit variance


- After applying batch normalization, we obtain the normalized  $\hat{h}_{l-1}$  that restores zero mean and unit variance
  - For almost any update to the lower layers,  $\hat{h}_{l-1}$  will remain a unit Gaussian
  - Output  $\hat{y}$  may be learned as a simple linear function  $\hat{y} = \mathbf{w}_l \hat{h}_{l-1}$
- Learning in this model is now very simple
  - Because parameters at the lower layers do not have an effect in most cases
- Their output is always renormalized to a unit Gaussian

# Batch Normalization -> learning easy

- Without normalization, updates would have an extreme effect of the statistics of  $h_{l-1}$
- Batch normalization has thus made this model easier to learn
- In this example the ease of learning came from making the lower layers useless
  - In our linear example: Lower layers not harmful but not beneficial either
    - Because we have normalized-out all effects of higher order than 1<sup>st</sup> and 2<sup>nd</sup> order stats
- In a deep neural network lower levels still can perform useful nonlinear transformations

# Reintroducing Expressive Power

- Normalizing the mean and standard deviation can reduce the expressive power of the neural network containing that unit
- To maintain the expressive power replace the batch of hidden unit activations  $\mathbf{H}'$  with  $\gamma\mathbf{H}' + \beta$ 
  - $\gamma$  and  $\beta$  are learned parameters that allow the new variable to have any mean and standard deviation
- Why did we normalize to zero and std. dev. when we then allow all means and deviations again?
  - Has different learning dynamics than unnormalized approach
  - It has the same power as  $\mathbf{H}$ , but the values of  $\mathbf{H}$  where determined by a complicated interaction between the parameters in the layers below
  - The new parametrization, the mean is only determined by  $\beta$



# Optimization for Deep Learning

- Parameter Initialization
- Batch Normalization
- **Pre-Training**

# Motivation

- Sometimes, directly training a model to solve a specific task can be too ambitious, if:
  - Model is too complex and hard to optimize
  - The task is very difficult
- It may be more effective to
  - Train a simpler model to solve the task, then move on to confront the final task
  - Methods collectively known as pretraining





# Greedy Supervised Pretraining

- Greedy Algorithm:
  1. Break a problem into many components
  2. Solve for the optimal version of each component in isolation
  3. Combine the solutions
- Combining the component solutions may not yield an optimal complete solution
- However, greedy algorithms can be computationally much cheaper than algorithms that solve for the best joint solution
- Quality of a greedy solution is often acceptable if not optimal
- Initializing the joint optimization algorithm with a greedy solution can speed it up and improve the quality of the solution

# Training each layer separately

- Supervised learning involving only a subset of the layers in the final neural network
- An example of greedy supervised pretraining
  - In which each added hidden layer is pretrained as part of a shallow supervised MLP
  - Taking as input the output of the previously trained hidden layer



# Extension to Transfer Learning

- Pretraining extends the idea to transfer learning
- Pretrain convolutional net with  $k$  layers on tasks
  - E.g. on a subset of 1000 ImageNet object categories
- Then initialize same-size network with the first  $k$  layers of the first net
  - All layers of second network (with upper layers initialized randomly) are then jointly trained to perform a different set of tasks
    - E.g. another subset of 1000 ImageNet categories, with fewer training examples than for the first set of tasks

# FitNets

- While depth improves performance, it also makes gradient-based training more difficult since deeper networks are more non-linear.
- Solution is to train a network with low enough depth (e.g. 5) and great enough width (no. of units per layer) to be easy to train
- This network becomes a teacher for a second network, designated the student
- Student network is much deeper and thinner and would be difficult to train with SGD, (e.g., 15-20 layers)

# Training the student network

- Task is made easier by training student network not only to predict output for original task, but also to predict value of middle layer of the teacher network
- This extra task provides a set of hints about how the hidden layers should be used and can simplify the optimization problem
- Additional parameters are introduced to regress the middle layer of the 5-layer teacher network from the middle layer of the deeper student network

# Predicting Intermediate Layers

- Instead of predicting the final classification target, the objective is to predict the middle hidden layer of the teacher network
- Objectives of Lower layers of student network:
  1. Help outputs of student network accomplish task
  2. Predict intermediate layer of the teacher network
- A thin-deep network may be more difficult to train than a wide-shallow network, but may generalize better and has lower computational cost if it is thin enough to have far fewer parameters.